

COMP6247: Reinforcement and Online Learning

Maheesan Niranjana

School of Electronics and Computer Science
University of Southampton

Sample of Case Studies

Spring Semester 2020/21

Review of Recent Papers

- **Li, K. and Malik, J. Learning to Optimize Neural Networks**, *arXiv preprint arXiv:1703.00441*
- Mnih, V. *et al.* Human-level control... [Atari games], *Nature* (2015)
- Yuan, X. *et al.* Reinforcement Learning for Elevator Control, *IFAC Conference* (2008)
- Popova, M. *et al.* Deep reinforcement learning for *de novo* drug design, *Science Advances* (2018)
- Mao, H. *et al.* Resource management with deep reinforcement learning, 15th *ACM Workshop on Hot Topics in Networks* (2016).

When reviewing for the assignment... look for:

- State and action spaces
 - Reward, cost function
 - Representations
 - Algorithm
 - Empirical work to illustrate the point
-
- Ignore some technical / algorithmic details outside the scope of what we have studied

I. Li and Malik: Learning to Optimize Neural Networks

- Training algorithm:
 - Objective function f
 - Initial conditions $\mathbf{x}^{(0)}$
 - Learning curve: $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(T)}$
- We do a lot of this:
 - f_1, f_2, \dots, f_n
- We tune hyper-parameters, choose step size, taper down learning rate...
- Can we learn from the experience of training networks, to efficiently train new networks?
- Reinforcement Learning formulation:
 - Step size selection as actions
 - Past updates as observation / state: o_t , features Ψ_t
 - Cost to minimize, rather than reward to maximize

Li and Malik: Learning to Optimize Neural Networks

Formulation

- Goal: Policy that minimizes *expected total* cost over time (epochs)

algorithm is to learn a policy π^* that minimizes the total expected cost over time. More precisely,

$$\pi^* = \arg \min_{\pi} \mathbb{E}_{s_0, a_0, s_1, \dots, s_T} \left[\sum_{t=0}^T c(s_t) \right],$$

where the expectation is taken with respect to the joint dis-

- Probability distribution over trajectory (state sequence)

$$q(s_0, a_0, s_1, \dots, s_T) = \int_{o_0, \dots, o_T} p_i(s_0) p_o(o_0 | s_0) \prod_{t=0}^{T-1} \pi(a_t | o_t, t) p(s_{t+1} | s_t, a_t) p_o(o_{t+1} | s_{t+1}).$$

- Policy: $\pi(a_t | o_t, t)$, conditional probability over actions
- $\pi(a_t | o_t, t) = \mathcal{N}(\mu^\pi(o_t), \Sigma^\pi(o_t))$
- $\mu^\pi(o_t)$ and $\Sigma^\pi(o_t)$ come from function approximators (neural networks)

Guided Policy Search

- Searching over large class of policies: π , probability distribution over actions (discrete / continuous)
- Idea is to maintain two policies:
 - ψ Linear, learned in closed form
 - π More expressive, desired
 - In each iteration, solve ψ and use it to train π .
- Objective function as constrained optimization problem:

More precisely, GPS solves the following constrained optimization problem:

$$\min_{\theta, \eta} \mathbb{E}_{\psi} \left[\sum_{t=0}^T c(s_t) \right] \text{ s.t. } \psi(a_t | s_t, t; \eta) = \pi(a_t | s_t; \theta) \quad \forall a_t, s_t, t$$

where η and θ denote the parameters of ψ and π respectively, $\mathbb{E}_{\rho}[\cdot]$ denotes the expectation taken with respect to the trajectory induced by a policy ρ and $\pi(a_t | s_t; \theta) := \int_{o_t} \pi(a_t | o_t; \theta) p_o(o_t | s_t)^2$.

Guided Policy Search (cont'd)

- Constrained optimization

$$\min_{\theta, \eta} E_{\psi} \left[\sum_{t=0}^T c(o_t) \right]$$

subject to $\psi(a_t | s_t, t; \eta) = \pi(a_t | s_t; \theta)$

- Relaxed version:

$$\min_{\theta, \eta} E_{\psi} \left[\sum_{t=0}^T c(o_t) \right]$$

subject to $E_{\psi}[a_t] = E_{\psi}[E_{\pi}[a_t | s_t]]$

- The simpler model and the desired model to be the same (select the same action) in expectation (on average).

Solution

Iterative algorithm (ADMM)

- Update groups of parameters

$$\eta \leftarrow \arg \min_{\eta} \sum_{t=0}^T \mathbb{E}_{\psi} [c(s_t) - \lambda_t^T a_t] + \nu_t D_t(\eta, \theta)$$

$$\theta \leftarrow \arg \min_{\theta} \sum_{t=0}^T \lambda_t^T \mathbb{E}_{\psi} [\mathbb{E}_{\pi} [a_t | s_t]] + \nu_t D_t(\theta, \eta)$$

$$\lambda_t \leftarrow \lambda_t + \alpha \nu_t (\mathbb{E}_{\psi} [\mathbb{E}_{\pi} [a_t | s_t]] - \mathbb{E}_{\psi} [a_t]) \quad \forall t,$$

where $D_t(\theta, \eta) := \mathbb{E}_{\psi} [D_{KL}(\pi(a_t | s_t; \theta) || \psi(a_t | s_t, t; \eta))]$
and $D_t(\eta, \theta) := \mathbb{E}_{\psi} [D_{KL}(\psi(a_t | s_t, t; \eta) || \pi(a_t | s_t; \theta))]$.

- $D_t(\theta, \eta) = E_{\psi} [D_{KL}(\pi(a_t | s_t; \theta) || \psi(a_t | s_t, t; \eta))]$
- Recall: $D_{KL}(p || q) = E_p [\log(p/q)] = \int_{-\infty}^{\infty} p(x) \log(p(x)/q(x)) dx$

- Setting the densities:

The algorithm assumes that $\psi(a_t | s_t, t; \eta) = \mathcal{N}(K_t s_t + k_t, G_t)$, where $\eta := (K_t, k_t, G_t)_{t=1}^T$ and $\pi(a_t | o_t; \theta) = \mathcal{N}(\mu_\omega^\pi(o_t), \Sigma^\pi)$, where $\theta := (\omega, \Sigma^\pi)$ and $\mu_\omega^\pi(\cdot)$ can be an arbitrary function that is typically modelled using a nonlinear function approximator like a neural net.

- Recall: linear transform; Multi-variate Gaussian!

$$\tilde{p}(s_{t+1} | s_t, a_t, t; \xi) = \mathcal{N}(A_t s_t + B_t a_t + c_t, F_t)$$

Draw samples of trajectory (using ψ) and fit linear model above.

- Quadratic approximation to cost function

\tilde{s} samples of s_t drawn from the trajectory induced by ψ so that $c(s_t) \approx \tilde{c}(s_t) := \frac{1}{2} s_t^T C_t s_t + d_t^T s_t + h_t$ for s_t 's that are near the samples.

With these assumptions, the subproblem that needs to be solved to update $\eta = (K_t, k_t, G_t)_{t=1}^T$ becomes:

$$\begin{aligned} \min_{\eta} \quad & \sum_{t=0}^T \mathbb{E}_{\tilde{s}} \left[\tilde{c}(s_t) - \lambda_t^T a_t \right] + \nu_t D_t(\eta, \theta) \\ \text{s.t.} \quad & \sum_{t=0}^T \mathbb{E}_{\tilde{s}} \left[D_{KL}(\psi(a_t | s_t, t; \eta) \| \psi(a_t | s_t, t; \eta')) \right] \leq \epsilon, \end{aligned}$$

- Heavy in detail **skip**; But, But, But... Have we seen this before?

- Recall Extended Kalman Filter:


- We had a nonlinear function
- We had quadratic cost
- We updated using local gradient
- Got a Gaussian approximation going from $n-1 | n-1$ to $n | n-1$

- Features (to define state)

Because of the stochasticity of gradients and objective values, the state features $\Phi(\cdot)$ are defined in terms of summary statistics of the history of iterates $\{x^{(i)}\}_{i=0}^t$, gradients $\{\nabla \hat{f}(x^{(i)})\}_{i=0}^t$ and objective values $\{\hat{f}(x^{(i)})\}_{i=0}^t$. We define the following statistics, which we will refer to as the average recent iterate, gradient and objective value respectively:

- $\overline{x^{(i)}} := \frac{1}{\min(i+1, 3)} \sum_{j=\max(i-2, 0)}^i x^{(j)}$
- $\overline{\nabla \hat{f}(x^{(i)})} := \frac{1}{\min(i+1, 3)} \sum_{j=\max(i-2, 0)}^i \nabla \hat{f}(x^{(j)})$
- $\overline{\hat{f}(x^{(i)})} := \frac{1}{\min(i+1, 3)} \sum_{j=\max(i-2, 0)}^i \hat{f}(x^{(j)})$

- Results



- Review:

- Show that you understand the framework
- Discuss if the results are persuasive
- Enjoy learning about the algorithm

II. Mnih, V. *et al.* Human-level control... [Atari games]

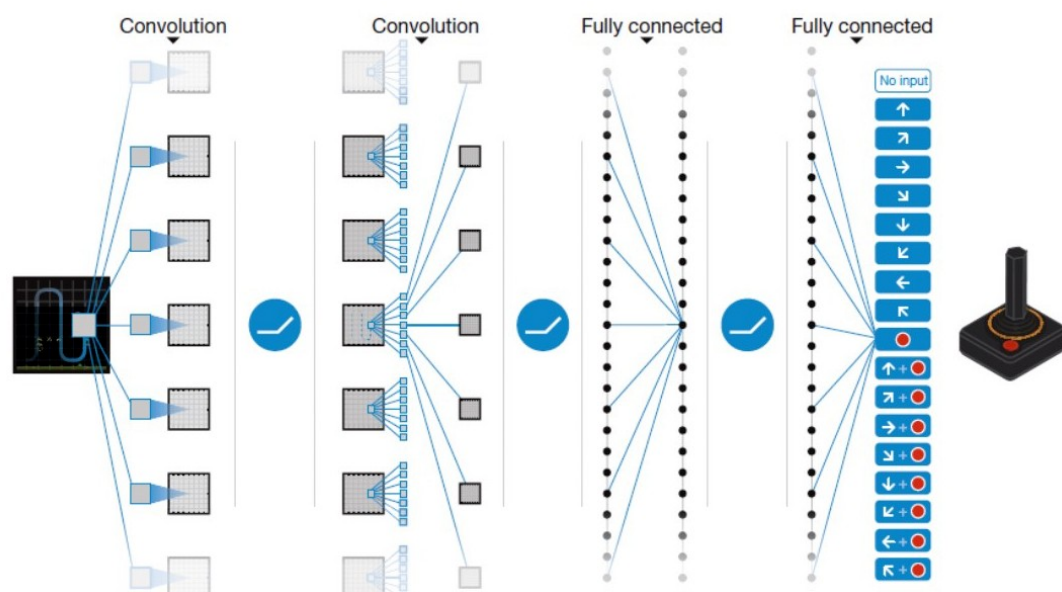
- Long history of computer games and machine learning
- Driven more by artificial intelligence than by statistical pattern recognition
- Sutton and Barto:
 - Samuel's checker player
 - Tesauro's Backgammon player
- Recent headline-grabbing developments in Atari Games (this paper), Go, Chess etc.
- What we are familiar with:

$$Q^*(s, a) = \max_{\pi} E \left[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = 1, \pi \right]$$

- $Q(s, a)$ represented by a function approximator (e.g. RBF in the mountain car problem)
- Temporal difference error between $Q(s, a)$ and $r + \gamma \max_{a'} Q(s', a')$
- Issues with scaling up; two solutions in this paper (**experience replay** and **updating frequency**)


II. Mnih, V. *et al.* Human-level control... [Atari games] (cont'd)

- Architecture




II. Mnih, V. *et al.* Human-level control... [Atari games] (cont'd)

Results



II. Mnih, V. *et al.* Human-level control... [Atari games] (cont'd)

Visualization



Experience Replay

II. Mnih, V. *et al.* Human-level control... [Atari games] (cont'd)

- Agent's experience at time t : $e_t = \{s_t, a_t, r_t, s_{t+1}\}$
- Dataset $D_t = \{e_1, e_2, \dots, e_t\}$
- Sample minibatches (subsets) from stored data: $U(D)$
- Minimize objective function:

$$L_i(\theta_i) = E_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

- Note:
 - iteration i , time t
 - θ_i^- Network parameters when we took the move
 - θ_i variable with respect to which we minimize

II. Mnih, V. *et al.* Human-level control... [Atari games] (cont'd)

When summarising:

- Note details in Supplementary material (style of journal)
- Core algorithmic setting
- How persuasive are results claimed
- What do they find by visualizing what the models learned (tSNE – similar to PCA, NMF)

III: Yuan, X. *et al.* Reinforcement Learning for Elevator Control

- Sum of rewards: $R_t = \sum_{i=0}^{\infty} \gamma^i r_{t+i+1}$
- Bellman: $Q^*(x, u) = \sum_{x' \in \mathcal{X}} f(x, u, x') [\rho(x, u, x') + \gamma \max_{u' \in \mathcal{U}} Q^*(x', u')]$
- Algorithms (good revision material):

Table 1. The Q-value iteration algorithm

```
Input  $f, \rho, \gamma$ , convergence threshold  $\theta$ 
Initialize  $Q_0$  arbitrarily, e.g.  $Q_0(x, u) = 0$ , for all  $x \in X, u \in U$ 
 $k = 0$ 
Repeat
  For each  $x \in X, u \in U$ 
     $Q_{k+1}(x, u) =$ 
      
$$\sum_{x' \in X} f(x, u, x') [\rho(x, u, x') + \gamma \max_{u' \in U} Q_k(x', u')]$$

  EndFor
   $k = k + 1$ 
Until  $\max_{x, u} |Q_k(x, u) - Q_{k-1}(x, u)| < \theta$ 
Output  $\pi^*(x) = \arg \max_{u \in U} Q_k(x, u) \quad \forall x \in X$ 
```

```
Input  $\gamma, \alpha_t$ , exploration parameters (e.g.,  $\varepsilon, \tau$ )
Initialize  $Q_0$  arbitrarily, e.g.  $Q_0(x, u) = 0$ , for all  $x \in X, u \in U$ 
Initialize  $x_0$ 
Repeat at each time step  $t$ :
  Choose  $u_t$  in  $x_t$  using policy derived from  $Q_t$ 
  (e.g. Boltzmann)
  Apply  $u_t$ , observe  $r_{t+1}, x_{t+1}$ 
   $Q_{t+1}(x_t, u_t) = Q_t(x_t, u_t) +$ 
    
$$\alpha_t [r_{t+1} + \gamma \max_{u \in U} Q_t(x_{t+1}, u) - Q_t(x_t, u_t)]$$

```

- Q Learning: $Q_{t+1}(x_t, u_t) = Q_t(x_t, u_t) + \alpha_t [r_{t+1} + \gamma \max_{u \in U} Q_t(x_{t+1}, u) - Q_t(x_t, u_t)]$
- Convergence; exploration

Elevator

• System Description

- The number of elevators (positive integer). In general, there are several elevators in elevator systems. In order to simplify the problem, we assume here that the system consists of a single elevator.
- The number of floors (positive integer). Set here to 5.
- The height of a floor (positive real). Set here to 6 m.
- The elevator speed (positive real). Here we set it to 3 m/s. This means the elevator takes 2 s to travel between two adjacent floors.
- The elevator capacity (positive integer). Set here to 4 passengers.
- The stop time, i.e., the sum of the intervals needed by the passengers to enter and exit the elevator on a floor. The stop time is set to 2 s. The fact that the

Elevator

What you should look for in summarising

- StateSpace

The state space of the elevator system is discrete and has 7 dimensions. The state signal x is composed of the following variables:

$$x = [c_1, c_2, c_3, c_4, p, v, o]^T. \quad (8)$$

where:

- c_i , $i = 1, 2, 3, 4$: Binary values, representing call requests (call flags) on each floor i . There is no call request on the ground floor in the down-peak traffic scenario.
- p : Discrete elevator position, taking values in $\{0, 1, 2, 3, 4\}$.
- v : Discrete vertical velocity taking values in $\{-3, 0, 3\}$ m/s.
- o : Discrete elevator occupancy, taking values in $\{0, 1, 2, 3, 4\}$. The number 0 means no passengers are inside the elevator; the number 4 means that the elevator is at its maximum capacity.

The cardinality of the state space is:

$$2^4 \cdot 5 \cdot 3 \cdot 5 = 1200 \quad (9)$$

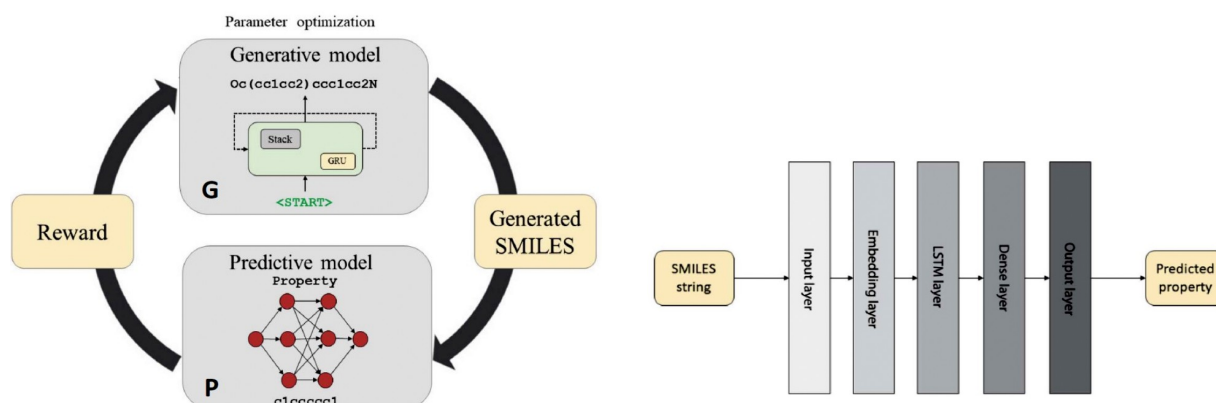
- Discrete action space; $\{-1, 0, 1\}$
- Constraints on actions
- Reward: $-\sum_{i=1}^4 c_i - 0$; (call requests and occupancy)
- Passenger arrival/departure models to simulate;
- Evaluation

IV: Popova, M. *et al.* Deep reinforcement learning for *de novo* drug design

- Biology, chemistry offer challenging machine learning problems
- Easy to construct molecules (synthetic chemistry)
- Difficult to experimentally validate their properties
- Learn from molecules whose properties are known; test new ones.
- Drug design: Molecular structure, representation, function prediction

Drug design (cont'd)

- A generative model and a predictive model!



Drug Design (cont'd)

What to look for in the paper

- States as representation of molecules as strings
- Terminal states when length T is reached
- Reward from the predictive model for a string of length T
- Policy $p(a_t|s_{t-1})$ enables the string to grow (the generative model)
- Policy network: $J(\Theta) = \sum_{s_T \in S^*} p_{\Theta}(s_T)r(s_T)$
- Solved by Policy Gradient methods (Section 13.1, Sutton and Barto)
 - $\pi(a|s, \theta)$, differentiable function of θ
 - REINFORCE class of algorithms

REINFORCE Algorithm

Monte Carlo Policy Gradient

- Cost function $J(\theta)$ and its gradient $\nabla J(\theta)$
- Stochastic gradient descent $\theta_{t+1} = \theta_t + \alpha, \widehat{\nabla J(\theta)}$
- Parameterize the policy (softmax):

$$\pi(a|s, \theta) = \frac{\exp(h(s, a, \theta))}{\sum_b \exp(h(s, b, \theta))}$$

- Policy gradient theorem ** skip **
- Algorithm: Eqn 13.7 & 13.8 of S & B.
- Monte Carlo: run multiple episodes and average (as seen before)

IV: Mao, H. *et al.* Resource management

- Captures ideas we discussed around the elevator and drug design problems.
- Optimize allocation of jobs to clusters in computing
 - Optimization problems
 - Gradient descent on (parameterized) policy
- See how the REINFORCE algorithm is specified in Eqn 2

$$\theta \leftarrow \theta + \alpha \sum_t \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) v_t$$

- Sample multiple trajectories and estimate reward v_t
- Look for simulation details and performance in review.