

COMP6247: Reinforcement and Online Learning

Maheesan Niranjana

School of Electronics and Computer Science
University of Southampton

Markov Decision Processes, Dynamic Programming

Chapters 3 and 4, Sutton and Barto

Spring Semester 2020/21


Foundations

- $S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3 \dots$
- Dynamics of the Markov Decision Process
 $p(s', r | s, a) = \Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$
- p is a function of four spaces: $\mathcal{S} \times \mathcal{R} \times \mathcal{A} \times \mathcal{S}$, mapping to $[0, 1]$
- $\sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} p(s', r | s, a) = 1 \quad \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ (Eqn 3.3 in S & B)
- This being a joint probability, we can extract other probabilities:
 - State transition: $p(s' | s, a) = \sum_{r \in \mathcal{R}} p(s', r | s, a)$
 - Expected Reward: $r(s, a) = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r | s, a)$
 - Expected reward for state - action - next state
 $p(r, a, s') = \sum_{r \in \mathcal{R}} \frac{p(s', r | s, a)}{p(s' | s, a)}$

Example: Recycling Robot

Page 52, Sutton and Barto


s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-



- Robot collects garbage cans
- Finite battery life, needs to recharge
- Search to collect or wait for can
- Reward when can collected

Example: Recycling Robot

State Transition Diagram



Example: Recycling Robot

States, Actions and Rewards

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-

- Homework: Read Example 3.3 in detail.

Returns and Episodes

- Of interest: expected return (not immediate reward)

$$G_t = R_{t+1} + R_{t+2} + \dots + R_T$$

- Assumed *episodes*: start to finish subsequences
- Each episode ends in *terminal state*
- Continuous tasks

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Discounting: $0 \leq \gamma \leq 1$
- Recursive Structure:

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots \\ &= R_{t+1} + \gamma (R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4} + \dots) \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned} \quad \text{Eqn.3.9, S\&B}$$

Value Functions

- Value function of a state ($v_\pi(s)$), under a policy π :

$$\begin{aligned}v_\pi(s) &= E_\pi [G_t | S_t = s] \\&= E_\pi \left[\sum_{k=1}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right]\end{aligned}$$

- Note: *total, expected*
- If you start from this state (s at time t) and execute policy π , what is your expected total reward?
- Similarly, value of action (at a state):



$$\begin{aligned}q_\pi(s, a) &= E_\pi [G_t | S_t = s, A_t = a] \\&= E_\pi \left[\sum_{k=1}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right]\end{aligned}$$

Recursive Structure and Bellman Equations

$$\begin{aligned}v(s) &= E_\pi [G_t | S_t = s] \\&= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s] \\&= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi [G_{t+1} | S_{t+1} = s']] \\&= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma v_\pi(s')] \quad \forall s\end{aligned}$$

- $v(s)$ written in terms of expectation over all actions and $v(s')$, the values of resulting states.
- Simultaneous equations with $v(s)$ as unknowns.

Gridworld Example



3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

- Rewards: -1 for going off grid; +10, +5 at A and B; 0 for all others
- Homework: Exercise 3.15, Page 61, S & B: “Adding a constant to all rewards does not change the relative values of states”.

Optimal Policies and Value Functions

- $v_*(s) = \max_{\pi} v_{\pi}(s)$
- Which policy gives us largest value at every state?
- Similarly: $q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$
- $q_*(s, a) = E[R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a]$

Bellman Equations for Optimal Value Functions


- For optimal state values

$$\begin{aligned}
 v_*(s) &= \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \\
 &= \max_a E_{\pi_*} [G_t | S_t = s, A_t = a] \\
 &= \max_a E_{\pi_*} [R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\
 &= \max_a E [R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \\
 &= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]
 \end{aligned}$$

- Similarly for optimal action values

$$\begin{aligned}
 q_*(s, a) &= E \left[R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') | S_t = s, A_t = a \right] \\
 &= \sum_{s', r} p(s', r | s, a) \left[r + \gamma \max_{a'} q_*(s', a') \right]
 \end{aligned}$$


Gridworld Example



Gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

v_*



π_*

Recycling Robot: Bellman Equations

Example 3.9, S & B

$$\begin{aligned} v_*(h) &= \max \left\{ \begin{array}{l} p(h|h, s)[r(h, s, h) + \gamma v_*(h)] + p(1|h, s)[r(h, s, 1) + \gamma v_*(1)], \\ p(h|h, w)[r(h, w, h) + \gamma v_*(h)] + p(1|h, w)[r(h, w, 1) + \gamma v_*(1)] \end{array} \right\} \\ &= \max \left\{ \begin{array}{l} \alpha[r_s + \gamma v_*(h)] + (1 - \alpha)[r_s + \gamma v_*(1)], \\ 1[r_w + \gamma v_*(h)] + 0[r_w + \gamma v_*(1)] \end{array} \right\} \\ &= \max \left\{ \begin{array}{l} r_s + \gamma[\alpha v_*(h) + (1 - \alpha)v_*(1)], \\ r_w + \gamma v_*(h) \end{array} \right\}. \end{aligned}$$

$$v_*(1) = \max \left\{ \begin{array}{l} \beta r_s - 3(1 - \beta) + \gamma[(1 - \beta)v_*(h) + \beta v_*(1)], \\ r_w + \gamma v_*(1), \\ \gamma v_*(h) \end{array} \right\}.$$

$$v_*(s) = \max_a \sum_{s', r} p(s', r|s, a) [r + \gamma v_*(s')]$$

Dynamic Programming

- Family of algorithms
- Policy Evaluation – State value function $v(s)$ for given policy π
- Bellman equations:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

- Turn the above into an iterative assignment *iterative policy evaluation*:

$$v_{k+1}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_k(s')]$$

- Variations in implementing updates
 - *expected* updates (after seeing all future states)
 - *sweeps* (in place) overwrite at every state visited

Iterative Policy Evaluation Algorithm

S & B, Page 75

- Evaluating a policy \rightarrow What is $v(s)$ under this policy?

Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input π , the policy to be evaluated

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Gridworld Example

S & B, Page 76




	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$
on all transitions

- $p(6, -1|5, \text{right}) = 1$, $p(7, -1|7, \text{right}) = 1$, $p(10, r|5, \text{right}) = 0$
- Undiscounted episodic tasks, terminal states at (shaded) corners.
- Iterate to improve policy

Gridworld Example



$$v_{k+1}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + v_k(s')]$$

Policy Iteration Algorithm

Page 80, S & B

Evaluate, improve...

$$\pi_0 \rightarrow V_{\pi_0} \rightarrow \pi_1 \rightarrow V_{\pi_1} \rightarrow \dots \pi_* \rightarrow V_*$$

- 1. Initialization**
 $V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$
- 2. Policy Evaluation**
 Repeat
 $\Delta \leftarrow 0$
 For each $s \in \mathcal{S}$:
 $v \leftarrow V(s)$
 $V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$
 $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
 until $\Delta < \theta$ (a small positive number)
- 3. Policy Improvement**
 $policy_stable \leftarrow true$
 For each $s \in \mathcal{S}$:
 $a \leftarrow \pi(s)$
 $\pi(s) \leftarrow \arg \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$
 If $a \neq \pi(s)$, then $policy_stable \leftarrow false$
 If $policy_stable$, then stop and return V and π ; else go to 2

Value Iteration Algorithm

Page 83, S & B

$$\begin{aligned}v_{k+1}(s) &= \max_a E[R_{t+1} + \gamma v_k(S_{t+1}) | S_t = s, A_t = a] \\&= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_k(s')]\end{aligned}$$

Initialize array V arbitrarily (e.g., $V(s) = 0$ for all $s \in \mathcal{S}^+$)

Repeat

$\Delta \leftarrow 0$

For each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number)

Output a deterministic policy, π , such that

$\pi(s) = \operatorname{argmax}_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

Summary

- Markov Decision Processes (MDP) in which we have knowledge of the environment
- Task is to evaluate a given policy and to find an optimal policy
- Solved by dynamic programming
- **Next: What if the environment is not known?**